

Represent tokens in context

AICE3002 / AICE6001

VLC = $\int \int \int$ Vision
Learning α VoLaC
Control

Tokenisation and Positional Representation

From discrete symbols to contextual representations

Jonathon Hare

Vision, Learning and Control
University of Southampton

Models operate on representations, not raw text

Words are an unstable unit of representation

Unknown words expose the vocabulary problem

Subwords balance coverage and sequence length

Byte-pair encoding builds a vocabulary by merging

Token IDs become learned vectors

Order must be represented explicitly

Fixed and learned positional representations

Position interacts with attention

Context changes what a token represents

From static embeddings to contextual representations