

**Attend
encode
classify**

AICE3002 / AICE6001

Transformer Encoders for Classification

Masking and a worked BERT-style classifier

Jonathon Hare

Vision, Learning and Control
University of Southampton

From self-attention to a Transformer block

Residual pathways and normalisation in the encoder

Multi-head attention preserves tensor shape

Padding masks prevent attention to absent tokens

Causal and bidirectional masks express different tasks

Stacking encoder blocks builds contextual representations

Pretraining creates a reusable encoder

The role of the classification token

Building a BERT-style classification head

Tracing shapes through the complete model

Fine-tuning the encoder and classifier together

Interpreting errors in a text classifier